

# **SWE016 Machine Learning Operations Project**

## **Deadline : 02.01.2026 Friday 12:00**

### **1. MLOps Term Project Announcement: Building and Governing Production ML Pipelines**

This project requires designing and implementing an end-to-end Machine Learning system demonstrating **MLOps Level 2 maturity: CI/CD Pipeline Automation**. Your final presentation must be a **business presentation** directed toward **senior management**, focusing on operational efficiency and risk mitigation.

#### **Term Project Theme: Developing a Resilient, High-Cardinality Prediction Service**

Your group must design, implement, deploy, and monitor an end-to-end ML system that tackles a supervised learning problem (classification or regression) involving complex or **high-cardinality categorical features**.

#### **I. Mandatory Deliverables and Rules**

Compliance with the following organizational and deliverable rules is mandatory:

Requirement	Details and MLOps Context
<b>Group Size</b>	Student groups must consist of <b>5 to 7 members</b> .
<b>Defined Roles</b>	Each student's <b>responsibilities must be clearly defined</b> . Map responsibilities to roles (e.g., Data Engineer, ML Engineer, DevOps).
<b>Working Demo</b>	The project <b>must include a working demo</b> . The demo should showcase the automated pipeline execution, the live prediction service, and the monitoring dashboard.
<b>Individual Report</b>	<b>Each student must upload their individual report</b> , detailing their contribution and technical implementation choices.
<b>Business Presentation (PPT)</b>	Prepare a PowerPoint presentation which should be a <b>business presentation</b> , assuming you will present this to <b>senior management</b> .
<b>Presentation Flow</b>	The presentation must strictly adhere to the <b>Introduction-Development-Conclusion flow</b> .
<b>Video Presentation</b>	Each group must prepare a <b>5-minute video presentation</b> at the end of the semester and share the link separately.
<b>In-Class Presentation</b>	The in-class presentation will be conducted <b>in person</b> , and the video presentation will also be required.

## **II. Required Tool Stack and MLOps Infrastructure**

Students are required to implement a modern, hybrid MLOps tool stack. The core tooling components are:

### **1. Experiment Tracking and Model Governance**

You may use **MLflow**.

- MLflow is the **dedicated experiment tracking tool** needed regardless of the chosen orchestrator.
- It must be used to log every parameter, metric, and model version.
- MLflow's **Model Registry** must be demonstrated as the central hub for version control and stage management (e.g., moving models from staging to production).

### **2. Workflow Orchestration and CI/CD Pipeline Engine (Choose One)**

Your group must select and justify one of the following tools to manage the overall pipeline workflow (DAG):

- **Kubeflow Pipelines (KFP):** Ideal if you focus on Kubernetes-native scale and deep learning.
- **Apache Airflow:** Ideal if focusing on reliable time-based automation for generalized data jobs and complex ETL (Data Preparation) steps.
- **Prefect:** Preferred for dynamic workflows that respond intelligently to data changes and fail fast during complex retraining loops.

### 3. Core CI/CD and Auxiliary Tools

The demonstration of CI/CD automation and resilient serving will require the use of standard DevOps and testing tools:

Tool Category	Example Tools Mentioned in Sources	Application in Project
Containerization	Docker, Kubernetes	Used to package the ML model, dependencies, and serve the Stateless Serving Function.
CI/CD Execution	Jenkins, GitLab CI/CD, Circle CI	Defining the pipeline steps (Commit Stage, Acceptance Test Stage).
Unit/Component Testing	xUnit Frameworks (JUnit, NUnit), DbUnit	Used for unit testing feature engineering logic and testing database interactions for repeatability.
Code Inspection/Analysis	Checkstyle, PMD, JDepend, NDepend	Ensuring code quality and architectural adherence in the CI pipeline.
Cloud	AWS, Azure, Google Cloud	All services are allowed

## III. Mandatory Technical Implementation Requirements

Your working demo must integrate and justify the use of specific Machine Learning Design Patterns, linking them to the chosen tool stack:

### 1. Data Representation (Handling High-Cardinality)

You must implement methods to transform complex inputs into features the model can operate on:

- **High Cardinality:** Implement and justify using either **EMBEDDINGS** (mapping high-cardinality data into a dense, lower-dimensional space, learning weights via gradient descent) or the **HASHED FEATURE** pattern (converting categorical input into a fixed number of buckets, accepting the trade-off of **bucket collision** and potential loss of model accuracy).
- **Feature Interactions:** Implement a **FEATURE CROSS** (concatenating two or more categorical features, potentially after **bucketing** numerical data) to allow the model to learn relationships between inputs faster.

## 2. Problem Representation and Training

Your solution must incorporate strategies for resilient modeling:

- **Problem Reframing:** Explore the use of **REFRAMING** (e.g., turning a regression into a classification problem by **bucketizing outputs**) to gain expressiveness of a full probability distribution or restrict the prediction range.
- **Performance Optimization:** Use the **ENSEMBLES** design pattern (such as Bagging or Boosting) to combine multiple models to decrease bias and/or variance, documenting the **Increased Training and Design Time** trade-offs.
- **Data Imbalance:** Apply the **REBALANCING** design pattern (e.g., Downsampling or Upsampling) to address highly skewed data, noting the importance of choosing an appropriate evaluation metric beyond simple accuracy.
- **Reproducibility:** Utilize **CHECKPOINTS** during training to ensure resilience (resuming training after interruption) and facilitate **fine-tuning** on fresh data by retraining from the last checkpoint.

## 3. Resilient Serving and Continuous Evaluation

The deployment must focus on reliability and monitoring:

- **Serving Architecture:** Implement the **STATELESS SERVING FUNCTION** pattern via a REST endpoint to handle prediction requests, noting its benefits for **autoscaling** and **language-neutral** operation. Discuss when the alternative, **BATCH SERVING**, would be more appropriate for asynchronous, non-latency-sensitive workloads.
- **Continuous Evaluation (CME):** Implement the **CONTINUED MODEL EVALUATION** design pattern to track model degradation caused by **data distribution shifts** (like concept drift or covariate shift).
  - Monitoring must track **ML-specific metrics** related to **features** and **predictions**.
  - The system must use evaluation metrics (like accuracy, precision, or AUC) to detect performance degradation.
  - Monitoring should incorporate statistical checks (like those available in tools such as **Great Expectations**) to perform feature validation and detect skew.
- **Algorithmic Fallback:** Present a strategy for **algorithmic fallbacks** (e.g., rolling back to an older model version or using a simpler, hard-coded solution) in case monitoring detects performance drops below acceptable thresholds.